




(a) Gradient norm vs. window length

(b) Parameter recovery error ($N = 10$ seeds)

(c) Training loss

(d) ρ convergence ($\rho_{\text{true}} = 28$, $\rho_{\text{init}} = 23$)